


PORTING LFD MINIAPP TO GPU VIA OPENMP OFFLOAD


Pankaj Rajak, Ye Luo and Aiichiro Nakano
ESP Project : NAQMC_RMD_ESP
Argonne National Laboratory

QXMD : Scalable Quantum Molecular Dynamics (QMD)



Zn porphyrin

- Open source program for QMD with capabilities for nonadiabatic QMD (NAQMD) and multiscale shock
- Follow the trajectory of all atoms while computing interatomic interaction from first principles in the framework of density functional theory (DFT)



Rubrene/C₆₀

quasi-electron;
quasi-hole

- *SoftwareX* 10, 100307: 1-5 (2019)
- *Proceedings of International Conference on High Performance Computing in Asia-Pacific Region, HPCAsia2020 best paper award, 1-10 (2020)*

LFD Miniapp for QXMD

- **Local field dynamics (LFD) : key computational kernel of NAQMD**
- **LFD solves many-electron dynamics in the framework of real-time (RT) time-dependent density functional theory**
- **input to LDF is potential field, and it return the electron density**

```
Total walltime      = 314.458 (s)
Electron-propagation = 92.8069 (s)
Field-propagation    = 208.392 (s)
calc_energy function = 26.7224 (s)
```

***Most expensive functions :
electron and field dynamics solvers***

- **Developed LFD mini-app in C++
for GPU offloading and
integration with QXMD**

Time spent in various functions on a test system

===== Time spent in initialization:

```
init_param      = 0.000279768 (s)
init_wfnvloc    = 0.228392 (s)
init_field      = 0.00852508 (s)
```

===== Time spent in individual functions:

```
kin_prop        = 89.1699 (s)
pot_prop        = 3.402 (s)
set_prop        = 0.0519754 (s)
compute_v       = 0.00362018 (s)
compute_vxc     = 0.563744 (s)
compute_rho     = 5.44145 (s)
periodic_bc     = 0.234753 (s)
spectral_field_solve = 0.0048161 (s)
```

Electron-propagation time ~ $\text{Kin_prop} + \text{pot_prop} + \text{periodic_bc}$

Electron filed solver: Kin_prop()

```
void kin_prop (int d, int p) {  
  float wrk[Nx+2][Ny+2][Nz+2][2], w[2];  
  for (int n=0; n < Norb; n++) {  
    for (int i=1; i <= Nr[0]; i++)  
      for (int j=1; j <= Nr[1]; j++)  
        for (int k=1; k <= Nr[2]; k++) {  
          w[0] = al[d][p][0]*psi[n][i][j][k][0] - al[d][p][1]*psi[n][i][j][k][1] ;  
          w[1] = al[d][p][0]*psi[n][i][j][k][1] + al[d][p][1]*psi[n][i][j][k][0] ;  
          ...  
          for (int s=0; s<2; s++) wrk[i][j][k][s] = w[s] ;  
        }  
    # update psi[n][i][j][k][s] ← wrk[i][j][k][s]  
  }  
}
```

- Inefficient Memory usage & loop structure
- By loop reordering, we can get rid of **wrk**
- **al** doesn't depend on n, j and k. (can be cached)

Electron filed solver: Kin_prop() Update-1

```
void kin_prop (int d, int p) {  
  float w[2];  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++)  
      for (int i=1; i <= Nr[0]; i++)  
        for (int n=0; n < Norb; n++) {  
          w[0] = al_0*psi[i][j][k][n][0] - al_1*psi[i][j][k][n][1] ;  
          w[1] = al_1*psi[i][j][k][n][1] + al_0*psi[i][j][k][n][0] ;  
          ...  
          # update psi[n][i][j][k][s] ← w[s]  
        }  
      }  
}
```

- Inefficient Memory usage & loop structure
- By loop reordering, we can get rid of **wrk**
- **al** doesn't depend on n, j and k. (can be cached)

- Better memory usage and data locality by changing data layout
 $\text{psi}[\mathbf{n}, i, j, k, s] \rightarrow \text{psi}[i, j, k, \mathbf{n}, s]$

Electron filed solver: Kin_prop() Update-2

```
void kin_prop (int d, int p) {  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++)  
      for (int i=1; i <= Nr[0]; i++)  
        for (int n=0; n < Norb; n++) {  
          w[0] = al_0*psi[i][j][k][n][0] - al_1*psi[i][j][k][n][1] ;  
          w[1] = al_1*psi[i][j][k][n][1] + al_0*psi[i][j][k][n][0] ;  
          w[0] += bl_0[i]*psi[i-i][j][k][n][0] - bl_1[i]*psi[i-i][j][k][n][1] ;  
          w[1] += bl_0[i]*psi[i-i][j][k][n][1] - bl_1[i]*psi[i-i][j][k][n][0] ;  
          ...  
          # update  $\psi[n][i][j][k][s] \leftarrow w[s]$   
        }  
      }  
  }  
}
```

➤ use old psi

Electron filed solver: Kin_prop() Update-2

```
void kin_prop (int d, int p) {  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++) {  
      for (int n=0; n <= Norb; n++) {  
        psi_old0[n] = psi[0][j][k][n][0];  
        psi_old1[n] = psi[0][j][k][n][1]; }  
      for (int i=1; i <= Nr[0]; i++)  
        for (int n=0; n < Norb; n++) {  
          w[0] = al_0*psi[i][j][k][n][0] - al_1*psi[i][j][k][n][1];  
          w[1] = al_1*psi[i][j][k][n][1] + al_0*psi[i][j][k][n][0];  
          w[0] += bl_0[i]*psi_old0[n] - bl_1[i]*psi_old1[n];  
          w[1] += bl_0[i]*psi_old1[n] - bl_1[i]*psi_old0[n];  
          ...  
          # update psi_old0 ← psi[n][i][j][k][0] and psi_old1 ← psi[n][i][j][k][s]  
          # update psi[n][i][j][k][s] ← w[s]  
        }  
      }  
    }  
}
```

➤ copy old psi to reduce scratch space

➤ Involves complex operation
➤ convert psi, psi_old into 1D complex variable

Electron filed solver: Kin_prop() Update-3

```
void kin_prop (int d, int p) {  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++) {  
      for (int n=0 ; n < Norb; n++)  
        psi_old[i] = psi[yz_stride+n];  
      for (int i=1; i <= Nr[0]; i++)  
        for (int n=0; n < Norb; n++) {  
          w = a1*psi[stride+n] + b1[i]*psi_old[n] + ...;  
          # update psi_old0[n] ← psi[stride+n]  
          # update psi[stride+n] ← w  
        }  
      }  
    }  
}
```

- Involves complex operation
- convert psi, psi_old into 1D complex variable

Before

```
std::float psi[Nx+2][Ny+2][Nz+2][Norb]  
std::float psi_old[Ny+2][Nz+2][Norb]
```

After


```
std::complex<float> psi  
std::complex<float> psi_old
```


Electron filed solver: Kin_prop() Offload

```
void kin_prop (int d, int p) {  
  #pragma omp team distribute collapse(2)  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++) {  
      #pragma omp parallel for simd nowait  
      for (int n=0 ; n < Norb; n++)  
        psi_old[i] = psi[yz_stride+n];  
      for (int i=1; i <= Nr[0]; i++)  
        #pragma omp parallel for simd nowait  
        for (int n=0; n < Norb; n++) {  
          w = a1*psi[stride+n] + b1[i]*psi_old[n] + ...;  
          # update psi_old0[n] ← psi[stride+n]  
          # update psi[stride+n] ← w  
        }  
    }  
}
```

Hierarchical parallelism

- Coarse grain parallelism via *omp team distribute* on outer loops
- Fine grain parallelism on inner Norb loop by *omp parallel for*
- Typical size of Nr is 256 and Norb 100



Electron filed solver: Kin_prop() Offload Timing

```
void kin_prop (int d, int p) {  
  #pragma omp team distribute collapse(2)  
  for (int j=1; j < Nr[1]; j++)  
    for (int k=1; k <= Nr[2]; k++) {  
      #pragma omp parallel for simd nowait  
      for (int n=0 ; n < Norb; n++)  
        psi_old[i] = psi[yz_stride+n];  
      for (int i=1; i <= Nr[0]; i++)  
        #pragma omp parallel for simd nowait  
        for (int n=0; n < Norb; n++) {  
          w = a1*psi[stride+n] + b1[i]*psi_old[n] + ...;  
          # update psi_old0[n] ← psi[stride+n]  
          # update psi[stride+n] ← w  
        }  
    }  
}
```

Updated timing

Total wall time	= 208.29 (s)
Electron-propagation time	= 1.44 (s)
Field-propagation time	= 206.08 (s)
calc_energy function time	= 18.81 (s)

Original

Total walltime	= 314.458 (s)
Electron-propagation	= 92.8069 (s)
Field-propagation	= 208.392 (s)
calc_energy function	= 26.7224 (s)

➤ *Using xlr compiler*

Field Dynamics Solver: Field_prop ()

```
void field_prop () {  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++)  
                vH[2][i][j][k] = fx*vH[0][i-1][j][k] +.... + rho*rho[i][j][k]  
    ...  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++){  
                vH[1][i][j][k] += vH[2][i][j][k];  
                vH[0][i][j][k] += vH[1][i][j][k]; }  
    ...  
}
```

- Contains multiple loops which updates the individual point of the 4D vH[3][Nx][Ny][Nz] grid

Field Dynamics Solver: Field_prop () Offload

```
void field_prop () {  
    #pragma omp target teams distribute parallel for simd collapse(3)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++)  
                vH[2*dim_stride+ stride] = fx*vH[offset_rho-xyz_stride] +...  
    ...  
    #pragma omp target teams distribute parallel for simd collapse(3)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++){  
                vH[1*dim_stride+ stride] += vH[2*dim_stride+ stride];  
                vH[0*dim_stride+ stride] += vH[1*dim_stride+ stride];  
            }  
    ...  
}
```

- Convert vH[3][Nx][Ny][Nz] into 1D dynamic array vH
- Allocate vH on device to minimize data movement between host and device
- Flat parallelism for omp offload

Field_prop () Offload Timing

```
void field_prop () {  
    #pragma omp target teams distribute parallel for simd collapse(3)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++)  
                vH[2*dim_stride+ stride] = fx*vH[offset_rho-xyz_stride] +...  
    ...  
    #pragma omp target teams distribute parallel for simd collapse(3)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++){  
                vH[1*dim_stride+ stride] += vH[2*dim_stride+ stride];  
                vH[0*dim_stride+ stride] += vH[1*dim_stride+ stride];  
            }  
    ...  
}
```

Updated timing

```
Total walltime      = 113.608 (s)  
Electron-propagation = 1.55114 (s)  
Field-propagation   = 111.271 (s) ←
```

Original

```
Total walltime      = 314.458 (s)  
Electron-propagation = 92.8069 (s)  
Field-propagation   = 208.392 (s)  
calc_energy function = 26.7224 (s) ←
```

GPU Activity time of field_prop

33.77%	9.84747s	1100000
19.38%	5.65139s	1100000
19.38%	5.65079s	1100000
19.34%	5.64034s	1100000

Field_prop () Asynchronous Offload

```
void field_prop () {  
    #pragma omp target teams distribute parallel for simd collapse(3) nowait depend(inout:vH_ptr)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++)  
                vH[2*dim_stride+ stride] = fx*vH[offset_rho-xyz_stride] +...  
    ...  
    #pragma omp target teams distribute parallel for simd collapse(3) nowait depend(inout:vH_ptr)  
    for (int i=1; i<=Nr[0]; i++)  
        for (int j=1; j<=Nr[1]; j++)  
            for (int k=1; k<=Nr[2]; k++){  
                vH[1*dim_stride+ stride] += vH[2*dim_stride+ stride];  
                vH[0*dim_stride+ stride] += vH[1*dim_stride+ stride];  
            }  
    ...  
}
```

Field_prop Timing

Original = 208.392 s

Sync. = 111.271 s

Async. = 26.722 s

Benchmark Results

System Size: $N_x=N_y=N_z=32$, $N_{orb}=64$, Unit-cell (1,2,1)

Branch	Electron-Propagation		Field-propagation		Total time	
	Xlr_c	ifx	Xlr_c	ifx	Xlr_c	ifx
Master	92.8069 s	8.841 s	208.39 s	41.88 s	314.45 s	51.52 s
Kin_offload	1.4400 s	2.06 s	206.08 s	28.03 s	208.29 s	31.17 s
Kin_Field_sync_offload	1.5511 s	Error	111.27 s	Error	113.60 s	Error
Kin_Field_async_offload	0.8317 s	-	27.46 s	-	29.03 s	-

Summary

1) System Size: $N_x=N_y=N_z=32$, $N_{orb}=32$, Unit-cell (1,1,1)

Branch	Electron-propagation (s)	Field-propagation (s)	Total Time (s)
Master	46.7649	42.0905	95.7997
Kin_offload	0.79	41.72	43.23
Kin_Field_sync_offload	0.752694	22.3408	23.8729
Kin_Field_async_offload	0.449227	5.56158	6.74293

2) System Size: $N_x=N_y=N_z=32$, $N_{orb}=64$, Unit-cell (1,2,1)

Branch	Electron-propagation (s)	Field-propagation (s)	Total Time (s)
Master	92.8069	208.392	314.458
Kin_offload	1.44	206.08	208.29
Kin_Field_sync_offload	1.55114	111.271	113.608
Kin_Field_async_offload	0.831752	27.457	29.0229