

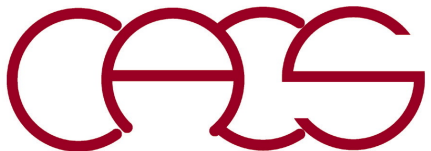
Split Molecular Dynamics

Aiichiro Nakano

*Collaboratory for Advanced Computing & Simulations
Department of Computer Science
Department of Physics & Astronomy
Department of Quantitative & Computational Biology
University of Southern California*

Email: anakano@usc.edu

Goal: Learn MPI communicator concept using *in situ* data analysis of molecular dynamics simulation



MPI_Comm_split()


```
MPI_Comm mothercomm, daughtercomm;  
int color, key;  
MPI_Comm_split(mothercomm, color, key, &daughtercomm);
```

- **MPI_Comm_split()** subdivides a communicator, **mothercomm**, into a set of daughter communicators, where processes of the same **color** belong to the same daughter communicator. Processes within each **color** are ranked according to **key**, or if **key** is the same, according to the rank in **mothercomm**. It returns a pointer to a daughter communicator, **daughtercomm**, to which the process belongs.
- **MPI_Comm_split()** is a simpler, higher-level function to construct communicators, instead of using **MPI_Comm_create()** combined with **MPI_Group_excl()** or **MPI_Group_incl()**.

MD & Analysis Communicators

- Split `MPI_COMM_WORLD` into two communicators; one performs molecular dynamics (MD) simulation, whereas the other analyzes simulation data on the fly in background.

```
int gid,sid,md;
MPI_Comm workers;
MPI_Comm_rank(MPI_COMM_WORLD,&gid); //Global rank
md = gid%2; // = 1 (MD workers) or 0 (analysis workers)
MPI_Comm_split(MPI_COMM_WORLD,md,0,&workers);
MPI_Comm_rank(workers,&sid); // Rank in workers
```



Run as `mpirun -n 2×nproc`

of processes needed for MD (specified in `pmd.h`)

Analysis: Velocity Probability Density

$P(v)$: Probability density function of atom velocity v

```
#define VMAX 5.0 // Maximum velocity value to construct a velocity histogram
#define NBIN 100 // # of bins in the histogram
```

```
void calc_pv() {
    double lpv[NBIN], pv[NBIN], dv, v;
    int i;
```

```
    dv = VMAX/NBIN; // Bin size
```

```
    for (i=0; i<NBIN; i++) lpv[i] = 0.0; // Reset local histogram
```

```
    for (i=0; i<n; i++) {
```

```
        v = sqrt(pow(rv[i][0],2)+pow(rv[i][1],2)+pow(rv[i][2],2));
```

```
        lpv[v/dv < NBIN ? (int)(v/dv) : NBIN-1] += 1.0; // Increment histogram
```

```
    }
```

```
    MPI_Allreduce(lpv, pv, NBIN, MPI_DOUBLE, MPI_SUM, workers);
```

```
    MPI_Allreduce(&n, &nglob, 1, MPI_INT, MPI_SUM, workers);
```

```
    for (i=0; i<NBIN; i++) pv[i] /= (dv*nglob); // Normalization
```


```
    if (sid == 0) {
```

```
        for (i=0; i<NBIN; i++) fprintf(fpv, "%le %le\n", i*dv, pv[i]);
```

```
        fprintf(fpv, "\n");
```

```
    }
```

```
}
```



$$v = |\vec{v}| = \sqrt{v_x^2 + v_y^2 + v_z^2}$$

$$\int dv p(v) \cong$$


$$dv \sum_{i=0}^{N_{\text{bin}}-1} p(v_i) = 1$$

Main Program: Initialization

```
init_params();  
if (md) {  
    set_topology();  
    init_conf();  
    atom_copy();  
    compute_accel();  
}  
else  
    if (sid == 0) fpv = fopen("pv.dat", "w");
```

Define FILE *fpv;

- All processes read input parameters, `init_params()`. The `nproc` processes of MD workers (`md == 1`) perform MD initialization tasks, whereas only rank 0 among the other `nproc` analysis workers (`md == 0`) opens a file to output the calculated velocity probability density function.



Main Program: Main MD Loop


```

for (stepCount=1; stepCount<=StepLimit; stepCount++) {
  if (md) single_step();
  if (stepCount%StepAvg == 0) {
    if (md) {
      Send # of atoms, n, to rank gid-1 in MPI_COMM_WORLD
      Send velocities of n atoms to rank gid-1 in MPI_COMM_WORLD
      eval_props();
    }
    else {
      Receive # of atoms, n, from rank gid+1 in MPI_COMM_WORLD
      Receive velocities of n atoms from rank gid+1 in MPI_COMM_WORLD
      calc_pv();
    }
  }
}

```

	0	1	2	3	4	5	
<i>dbuf</i>	v_{0x}	v_{0y}	v_{0z}	v_{1x}	v_{1y}	v_{1z}	...

$dbuf[3*i+a] \leftarrow rv[i][a] \ (i = 0, \dots, n-1; a = 0,1,2)$




- MD workers perform MD simulation. Every `stepAvg` steps, MD workers send their atom velocities to corresponding analysis workers (*i.e.*, those with the same ranks in respective daughter communicators). Upon receiving the velocities, analysis workers calculate the velocity probability density function.


Main Program: Finalization

```
if (md && sid == 0)
    printf("CPU & COMT = %le %le\n",cpu,comt);
if (!md && sid == 0)
    fclose(fpv);
```


- Rank 0 of MD workers reports the computing & communication times, whereas rank 0 of analysis workers closes the probability density output file.
- **Finally:** Change all MPI_COMM_WORLD's in the original MD functions to workers ~ it's only a matter in the small MD world!




Results



```
pmd.in
5 5 5 // InitUcell[0|1|2]
0.8 // Density
1.0 // InitTemp
0.005 // DeltaT
30 // StepLimit
10 // StepAvg
```



Maxwell-Boltzmann Distribution

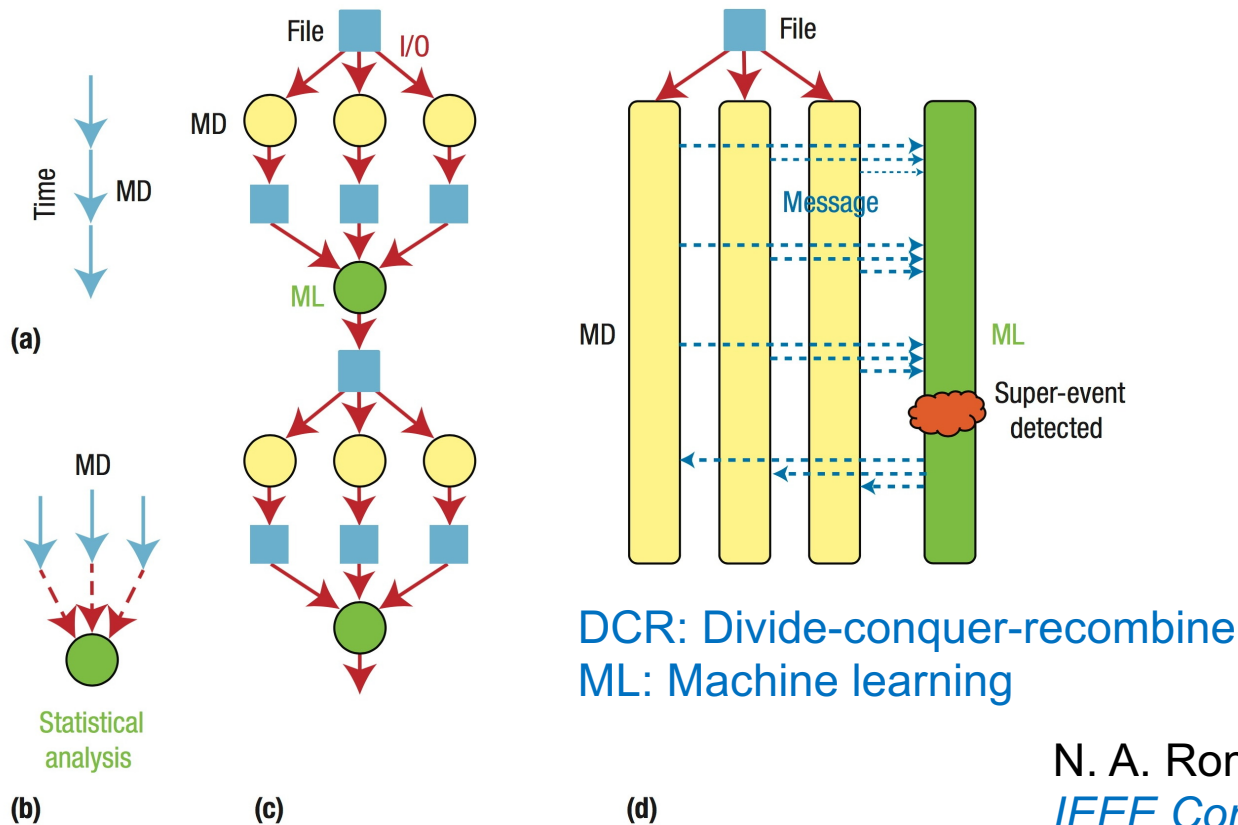


$$P_{\text{Maxwell-Boltzmann}}(v) = \frac{4}{\sqrt{\pi}} \left(\frac{m}{2k_{\text{B}}T} \right)^{3/2} v^2 \exp\left(-\frac{mv^2}{2k_{\text{B}}T}\right)$$

K. Shimamura *et al.*, [Appl. Phys. Lett. 107, 231903 \('15\)](#)

In Situ Data Analysis

Use communicators to add data analytics & extra logic to parallel simulations



N. A. Romero *et al.*,
[*IEEE Computer* **48**\(11\), 33 \('15\)](#)

FIGURE 2. DCR in time. (a) Molecular dynamics (MD) simulations have sequential time dependence. (b) Parallel replica dynamics (PRD) predicts long-time behavior through statistical analysis of multiple parallel MD trajectories. (c) Conventional file-based and (d) new in situ PRD simulations. ML represents machine-learning tasks.

See also T. Do *et al.*, [A lightweight method for evaluating *in situ* workflow efficiency](#), *J. Comput. Sci.* **48**, 101259 ('21)