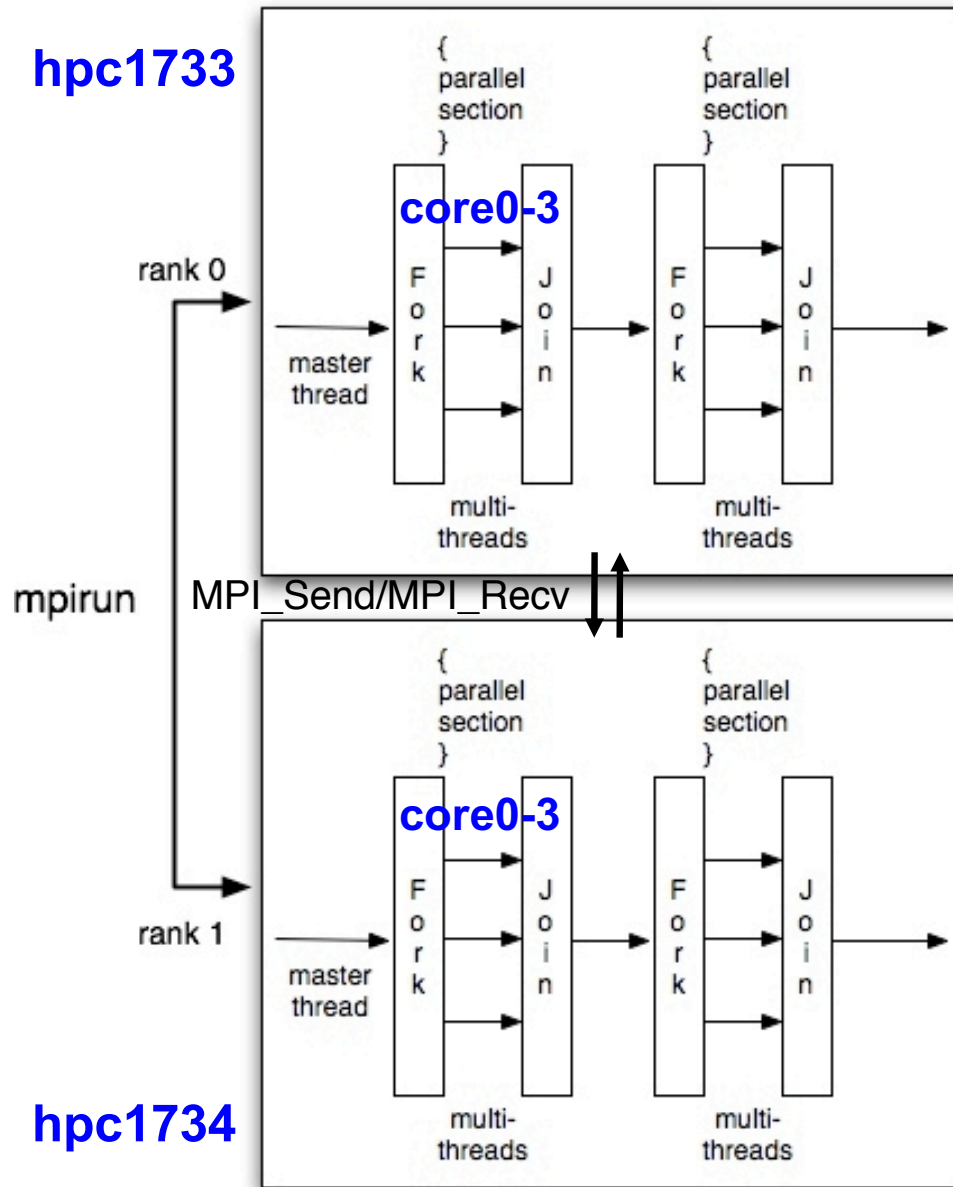


How Hybrid MPI+OpenMP MD Runs



In hmd.h:

```
int vproc[3] = {1,1,2}, nproc = 2;  
int vthrd[3] = {2,2,1}, nthrd = 4;
```

In hmd.c:

```
omp_set_num_threads(nthrd);
```

On hpc-login3:

```
salloc --nodes=2 --ntasks-per-node=1  
--cpus-per-task=4 -t 30
```

On hpc1733:

```
srun -n 2 ./hmd
```

On hpc1733 & hpc1734:


```
top (then type H and 1)
```

More on Multithreading MD

- Large overhead is involved in opening an OpenMP parallel section
→ Open it only once in the main function

In hmdm.c:

```
int main() {  
    ...  
    omp_set_num_threads(nthrd);  
    #pragma omp parallel  
    {  
        #pragma omp master  
        { // Do serial computations here }  
        ...  
        #pragma omp barrier // When threads need be synchronized  
        ...  
    }  
    ...  
}
```



More on Avoiding Race Conditions

- Program `hmd.c`: (1) used data privatization; (2) disabled the use of Newton's third law → this doubled computation
- **Cell-coloring**
 - > Race condition-free multithreading without duplicating pair computations
 - > Color cells such that no cells of the same color are adjacent to each other
 - > Threads process cells of the same color at a time in a color loop

1	3	1	3	1	3
0	2	0	2	0	2
1	3	1	3	1	3
0	2	0	2	0	2
1	3	1	3	1	3
0	2	0	2	0	2

H. S. Byun *et al.*,
Comput. Phys. Commun.
219, 246 ('17)

- Use graph coloring in more general computations

False Sharing

- While eliminating race conditions by data privatization, the use of consecutive per-thread accumulators, `lpe_td[nthrd]`, degrades performance by causing excessive cache misses

See [false sharing](#) Wiki page

- **Solution 1: Padding**


```
struct lpe_t {  
    double lpe;  
    double pads[7]; // assume intel CPU with 64 byte cache line  
};  
struct lpe_t lpe_td[nthrd];
```

- **Solution 2: System-supported data privatization**

```
#pragma omp parallel private (...) reduction(+:lpe)  
{  
    ...  
    lpe += 0.5*vVal;  
    ...  
}  
// No reduction over the threads is required here
```

1. Create private copies of the variable (`lpe`) in the reduction clause for all the threads
2. Perform the specified reduction operation (+) on the variable at the end of the parallel section

Scalability Test



Some Like It as Arguments

- Use command line arguments for scaling tests without re-compiling multiple times
- `hmd.c` → `hmdarg.c` by adding the following lines in `main()`

```
int main(int argc, char **argv) {  
    ...  
    vthrd[0] = atoi(argv[1]);  
    vthrd[1] = atoi(argv[2]);  
    vthrd[2] = atoi(argv[3]);  
    nthrd = vthrd[0]*vthrd[1]*vthrd[2];  
    printf("Number of threads = %d\n", nthrd);  
}
```

string-to-integer conversion (pointing to `atoi(argv[1])`)

command-line argument (pointing to `argv[2]`)

- **Compiling**

```
mpicc -o hmdarg hmdarg.c -fopenmp -lm
```

Strong-Scaling Test with hmdarg.c

```
[anakano@hpc-login3 cs596]$ salloc --nodes=1 --ntask-per-node=1 --cpus-per-
task=8 -t 59
...
[anakano@hpc1727 cs596]$ srun -n 1 ./hmdarg 1 1 1
Number of threads = 1
a1 = 4.103942e+01 4.103942e+01 4.103942e+01
lc   = 16 16 16
rc   = 2.564964e+00 2.564964e+00 2.564964e+00
thbk = 16 16 16
nglob = 55296
CPU & COMT = 1.073547e+01 2.005649e-02
[anakano@hpc1727 cs596]$ srun -n 1 ./hmdarg 2 1 1
Number of threads = 2
...
thbk = 8 16 16
nglob = 55296
CPU & COMT = 6.804797e+00 1.980424e-02
[anakano@hpc1727 cs596]$ srun -n 1 ./hmdarg 2 2 1
Number of threads = 4
...
thbk = 8 8 16
nglob = 55296
CPU & COMT = 4.956142e+00 1.981378e-02
[anakano@hpc1727 cs596]$ srun -n 1 ./hmdarg 2 2 2
Number of threads = 8
...
thbk = 8 8 8
nglob = 55296
CPU & COMT = 4.078273e+00 2.253795e-02
```